



Introduction to Hibernate

Object/Relational Mapping


- ▶ A major part of any enterprise application development project is the persistence layer
 - ▶ Accessing and manipulate persistent data typically with relational database
- ▶ ORM handles Object-relational impedance mismatch
 - ▶ Data lives in the relational database, which is table driven (with rows and columns)
- ▶ Relational database is designed for fast query operation of table-driven data
 - ▶ Work with objects, not rows and columns of table

Object-Relational Mapping

- ▶ Automated persistence of object to tables in RDBMS.
- ▶ Usually with the help of **metadata** that describes the mapping.
 - ▶ SQL is auto-generated by the metadata description.
- ▶ An ORM Solution consists of the following pieces:
 - ▶ Persistence Manager with CRUD API.
 - ▶ Query API
 - ▶ Mapping metadata

Domain Model and the paradigm mismatch

- ▶ Classes implement the business entities of our domain model
 - ▶ attributes of the entity are properties of our Java class
 - ▶ associations between entities are also implemented with properties



- ▶ Let's see if there is a problem mapping this to tables and columns...

Creating tables for the Domain Model

```
create table USER (  
    USER_NAME      varchar not null primary key,  
    ADDRESS         varchar not null)  
  
create table BILLING_DETAILS (  
    ACCOUNT_NUMBER  varchar not null primary key,  
    ACCOUNT_NAME    varchar not null,  
    ACCOUNT_TYPE    varchar not null,  
    USER_NAME       varchar foreign key references USER)
```

- We'll see the 5 problems of the O/R paradigm mismatch appear as we gradually make our model more complex...

The problem of granularity

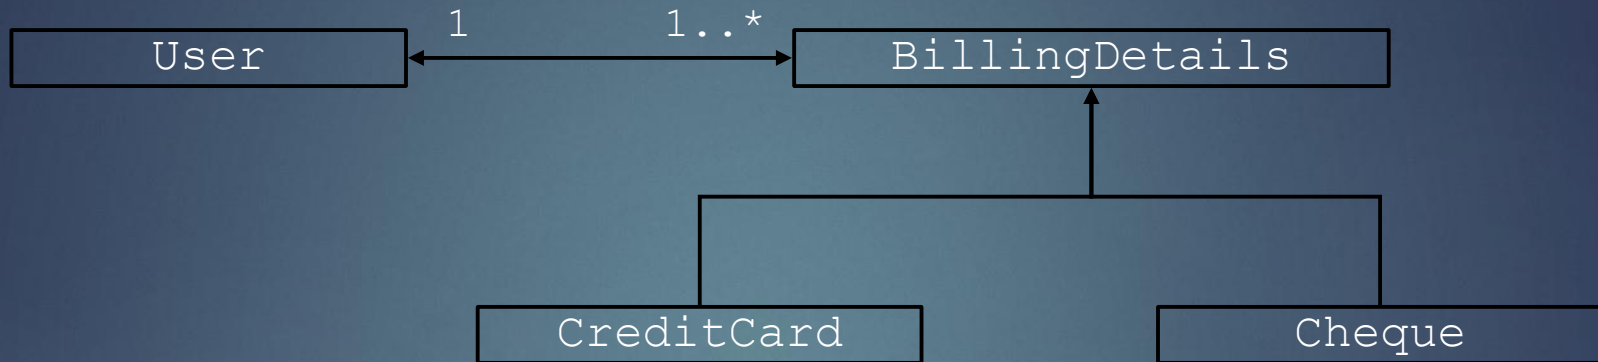


- ▶ should we create a new ADDRESS table?
 - ▶ should we create a new SQL data type and change the column?
 - ▶ user-defined data types (UDT) are not portable and the standard is weak
- ▶ We usually add new columns to USER with built-in SQL data types:

```
create table USER (  
    USER_NAME          varchar not null primary key,  
    ADDRESS_STREET     varchar not null,  
    ADDRESS_CITY       varchar not null,  
    ADDRESS_ZIPCODE    varchar not null)
```

The problem of subtypes

We create subclasses of BillingDetails:



and use polymorphism in Java to implement our billing strategy.

How do we represent subtypes in our relational model?

The problem of identity

In Java, we have two notions of "sameness"

- ▶ object identity is the memory location of an object, `a==b`
- ▶ object equality (what is this really?), `a.equals(b)`

In SQL databases, we identify a particular row using the primary key and the table name.

The problem of associations

- ▶ Object-oriented languages represent entity relationships as
 - ▶ *object references* (pointers) and *collections* of object references
- ▶ Relational databases represent entity relationships as
 - ▶ copies of primary key values
 - ▶ referential integrity ensured by *foreign key* constraints
- ▶ The mismatch:
 - ▶ object references are directional, there is no such concept in the relational model
 - ▶ many-to-many associations require a *link table* in relational databases

The problem of object graph navigation

In Java, we "walk" the object graph by following references:

```
 david.getBillingDetails().getAccountName()
```

In SQL, we join tables to get the required data:

```
select * from USER u
  left outer join BILLING_DETAILS bd
    on bd.USER_ID = u.USER_ID
 where u.USERNAME = "david"
```

The cost of the mismatch

These problems can, at least theoretically, be solved using handwritten SQL/JDBC

- ▶ by writing a lot of tedious code
- ▶ The “mismatch problem” is real
- ▶ better UDT support in SQL will not solve all the issues
- ▶ not all applications are suitable for table-oriented approaches

Is the solution design patterns (DAO)
or programming models (EJB entity beans)?

"How should we implement the persistence layer in our
application?"

Implement a Persistence Layer

- ▶ **EJB 2.0 entity beans**

- ▶ Bean Managed Persistence
- ▶ Container managed persitence
- ▶ No polymorphic associations and queries
- ▶ Not portable between different application servers
- ▶ Forces an unnatural Java Style

- ▶ **Object Oriented Database Systems**

- ▶ Not very popular

- ▶ **XML persistence**

Overview of Hibernate

- ▶ Object/relational mapping framework for enabling transparent POJO persistence
 - Open Source
- ▶ Persistence for JavaBeans, Model is not tied to persistence Implementation
- ▶ Powerful and Sophisticated queries (Criteria and HQL)
- ▶ Allows developers focus on domain object modeling not the persistence plumbing
- ▶ Performance
 - High performance object caching
 - Configurable materialization strategies
- ▶ Does not require a container

Persistent Objects and Collections


- ▶ Short-lived, single threaded objects containing persistent state and business function
- ▶ These might be ordinary JavaBeans/POJOs,
 - ▶ Not associated with (exactly one) Session
- ▶ Changes made to persistent objects are reflected to the database tables (when they are committed)
- ▶ As soon as the Session is closed, they will be detached and free to use in any application layer

Entity


- ▶ Lightweight persistent domain object – that which is persisted
- ▶ Restrictions
 - ▶ must have a public or protected no-arg constructor
 - ▶ cannot be final
 - ▶ cannot have final methods or final instance variables that are to be persisted
 - ▶ can be abstract or concrete class
 - ▶ must have a primary key

What is Persistence ?

- ▶ Ability of an object to survive even current session or program terminate.
- ▶ The ability of an object to remain in existence past the lifetime of the program that creates it.
- ▶ Its Achieved by
 - ▶ Relational Databases
 - ▶ Serialization
 - ▶ EJB (entity beans)



Hibernate Architecture



Hibernate Core Interfaces

Session Factory Interface & Classes

- ▶ Used to create Session objects, created during application initialization, caches generated SQL and other metadata

▶ Session Interface

- ▶ Primary interface, used to store & retrieve objects

▶ Configuration Class

- ▶ Used to configure & bootstrap Hibernate

▶ Transaction Interface

- ▶ May not be used by applications

▶ Query & Criteria Interface

- ▶ Query interface is used to run queries in HQL or SQL
- ▶ Criteria interface is also very similar

SessionFactory

- ▶ Represented by [org.hibernate.SessionFactory](#)
- ▶ A factory for *Session* and a client of **ConnectionProvider**
 - ▶ Typically one for each database
 - ▶ Maintains a threadsafe (immutable) cache of compiled mappings for a single database
 - ▶ Might hold an optional (second-level) cache of data that is reusable between transactions, at a processor cluster-level

Session

▶ Session

- ▶ Represented by [org.hibernate.Session](#)
- ▶ The life of a Session is bounded by the beginning and end of a logical transaction.
- ▶ A session represents a persistence context
- ▶ Handles life-cycle operations- create, read and delete operations - of persistent objects
- ▶ A single-threaded, short-lived object representing a conversation between the application and the persistent store
- ▶ Wraps a JDBC connection
- ▶ Factory for *Transaction*

Configuration

- ▶ Represented by `org.hibernate.config.Configuration`
- ▶ Configurations done using **.properties** and **XML** Files
- ▶ Hibernate is part of the app, and so is responsible for getting connections to the database
- ▶ Configuration file tells Hibernate how to get the connections
- ▶ Hibernate creates and manages a connection pool

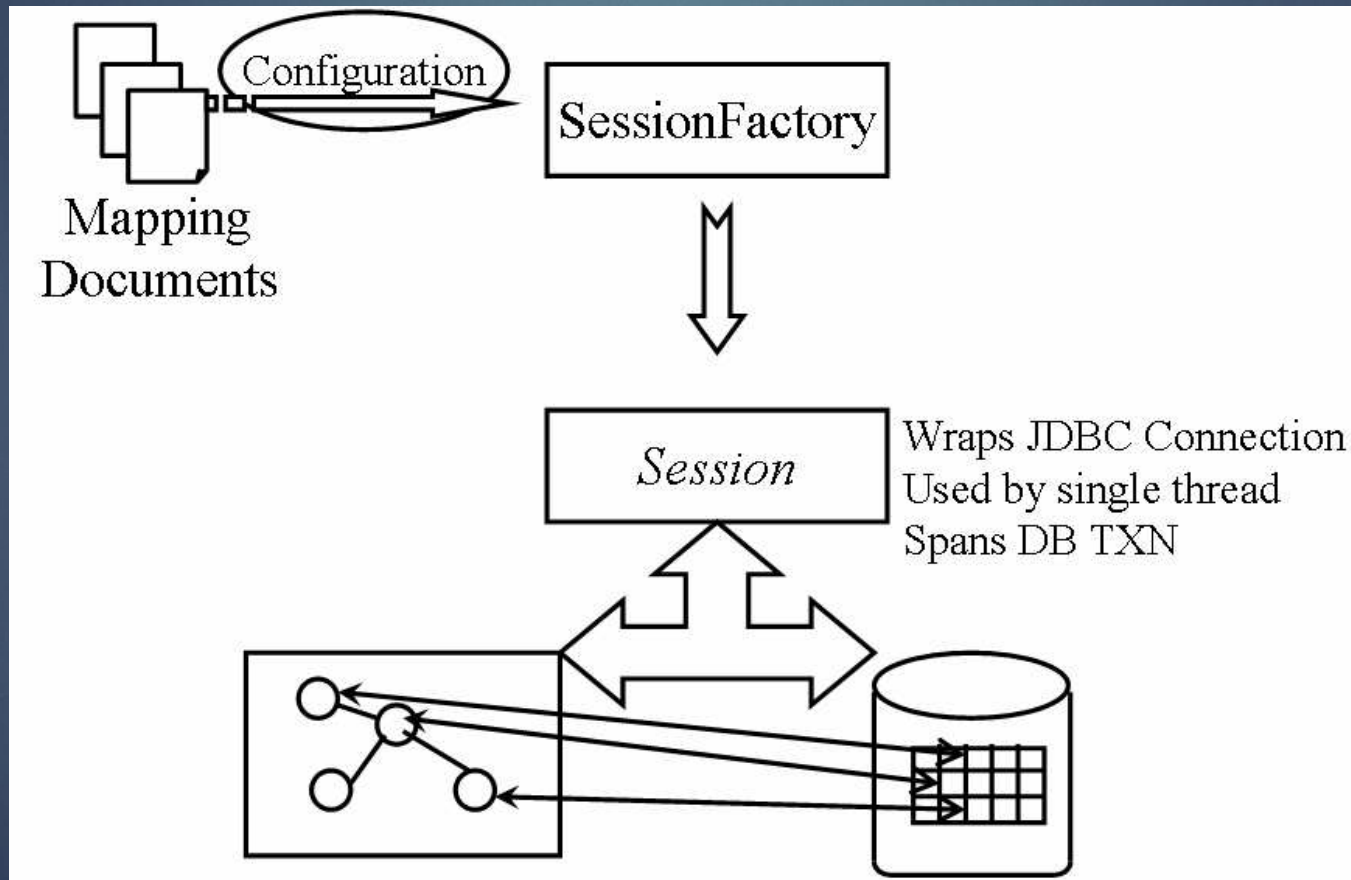
Configuring Hibernate

- ▶ Programmatic configuration
- ▶ XML configuration file
 - ▶ Specify a full configuration in a file named *hibernate.cfg.xml*
 - ▶ Configure the class to table mapping in **.hbm.xml*
 - ▶ By default, is expected to be in the root of your classpath
- ▶ Annotation Based Configuration

Hibernate Mapping

Hibernate uses runtime reflection to determine persistent properties of classes.

A mapping property or configuration file is used to generate database schema and provide persistence



Hibernate.cfg.xml

<!-- DTD →

```
<hibernate-configuration>
```

```
<session-factory name="myfact">
```

```
<property
name="hibernate.connection.driver_class">com.mysql.jdbc.Driver</property
>
```

```
<property name="hibernate.connection.password">srivatsan</property>
```

```
<property
name="hibernate.connection.url">jdbc:mysql://localhost:3306/employee</p
roperty>
```

```
<property name="hibernate.connection.username">root</property>
```

```
<property name="show_sql">>true</property>
```

```
<property
name="hibernate.dialect">org.hibernate.dialect.MySQLDialect</property>
```

```
<property name="hibernate.hbm2ddl.auto">update</property>
```

```
<mapping resource="Invoice.hbm.xml"/>
```

```
</session-factory>
```

```
</hibernate-configuration>
```

Dialects

- ▶ MySQL
 - ▶ **org.hibernate.dialect.MySQLDialect**
- ▶ Oracle
 - ▶ **org.hibernate.dialect.OracleDialect**
 - ▶ **org.hibernate.dialect.Oracle9Dialect**
- ▶ Microsoft SQL Driver
 - ▶ **org.hibernate.dialect.SQLServerDialect**
- ▶ Postgres
 - ▶ **org.hibernate.dialect.PostgreSQLDialect**

Domain Classes

- ▶ Domain classes are classes in an application that implement the entities of the business domain (e.g. Customer and Order in an Ecommerce application)
- ▶ Hibernate works best if these classes follow some simple rules, also known as the Plain Old Java Object (POJO) programming model.
- ▶ Hibernate assumes very little about the nature of your domain classes
 - ▶ You may express a domain model in other ways: using trees of Map instances, for example.

Steps to write a Domain Class

▶ Step 1: Implement a no-argument Constructor

- ▶ All persistent classes must have a default constructor so that Hibernate can instantiate them using

▶ Step 2: Provide an identifier property

- ▶ This property maps to the primary key column of a database table.
- ▶ The property might have been called anything, and its type might be any primitive type, any primitive "wrapper" type, *java.lang.String* or *java.util.Date*
- ▶ Composite key is possible

▶ Step 3: Declare accessor methods for persistent fields

- ▶ Hibernate persists JavaBeans style properties

The POJO Class

```
public class Invoice {
```

```
    private int invno;  
    private String customerName;  
    private double amount;
```

```
    public Invoice() {  
        super();  
    }
```

```
    public int getInvno() {  
        return invno;  
    }  
    public void setInvno(int invno) {  
        this.invno = invno;  
    }  
    public String getCustomerName() {  
        return customerName;  
        //OtherSet/Get Methods  
    }
```

Identifier property

Constructor –Zero Arg

Accessor/Mutator
Methods

Saving Objects

- ▶ To persist the object to the database when save is called with a valid Hibernate session
- ▶ An object remains to be in “transient” state until it is saved and moved into “persistent” state
- ▶ The class of the object that is being saved must have a mapping file (*invoice.hbm.xml*)

Getting Session Factory & Saving

```
public static void main(String[] args) {  
  
    SessionFactory fact =  
        new Configuration().configure().buildSessionFactory();  
  
    Session session=fact.openSession();  
  
    Transaction tx = session.beginTransaction();  
  
    Customer cust =new Customer(101,"Ramesh",4040);  
  
    session.save(cust);  
  
    tx.commit();  
  
}
```

Loading Objects

- ▶ Used for loading objects from the database
- ▶ Each *load(..)* method requires object's primary key as an identifier
 - ▶ The identifier must be *Serializable* – any primitive identifier must be converted to object
- ▶ Each *load(..)* method also requires which domain class or entity name to use to find the object with the id
- ▶ When Object exist use *load()*, - will throw an exception if the unique id is not found in the database
- ▶ *get()* method ,will return null if the unique id is not found in the database
- ▶ The *load()* method may return a proxy instead of a real persistent instance.
- ▶ On the other hand, *get()* never returns a proxy.

Loading objects

► From Session interface

► Session ses = sfact.openSession();

Person persObj =

(Person)sfact.openSession().get(Person.class, 10);

System.out.println("Name"+persObj.getName());

Person person = (Person) session.get(Person.class, id);

if (person == null){

System.out.println("Person is not found for id " + id);

ses.close();

Life-cycle Operations and SQL commands

- ▶ *save()* result in an *SQL INSERT*
- ▶ *delete()* results in an *SQL DELETE*
- ▶ *update()* results in an *SQL UPDATE*
- ▶ Changes to persistent instances are detected at flush time and also result in an *SQL UPDATE*

Updating Objects

- ▶ Hibernate automatically manages any changes made to the persistent objects
- ▶ The objects should be in “persistent” state not transient state
- ▶ If a property changes on a persistent object, Hibernate session will perform the change in the database when a transaction is committed (possibly by queuing the changes first)
- ▶ From developer perspective, you do not have to any work to store these changes to the database
- ▶ You can force Hibernate to commit all of its changes using *flush()* method

Updating Objects

- ▶ Get a Persistent Object by Calling Load
- ▶ Set a new Value for the field that need to be updated
- ▶ Call the update Method

```
Invoice inv = (Invoice)sess.load(Invoice.class, new  
Integer(idno));
```

```
inv.setCustomerName("Ganesh Kumar");
```

```
sess.update(inv);
```

```
sess.getTransaction().commit();
```


Deleting Objects

- ▶ Remove a persistent instance from the datastore.
- ▶ The argument may be an instance associated with the calling Session or a transient instance with an identifier associated with existing persistent state.
- ▶ This operation cascades to associated instances if the association is mapped with cascade="delete".

```
Invoice inv = (Invoice)sess.load(Invoice.class, new Integer(idno));
```

```
sess.delete(inv);
```

```
sess.getTransaction().commit();
```

```
System.out.println("Deleted");
```


Hibernate Life Cycle

transient state

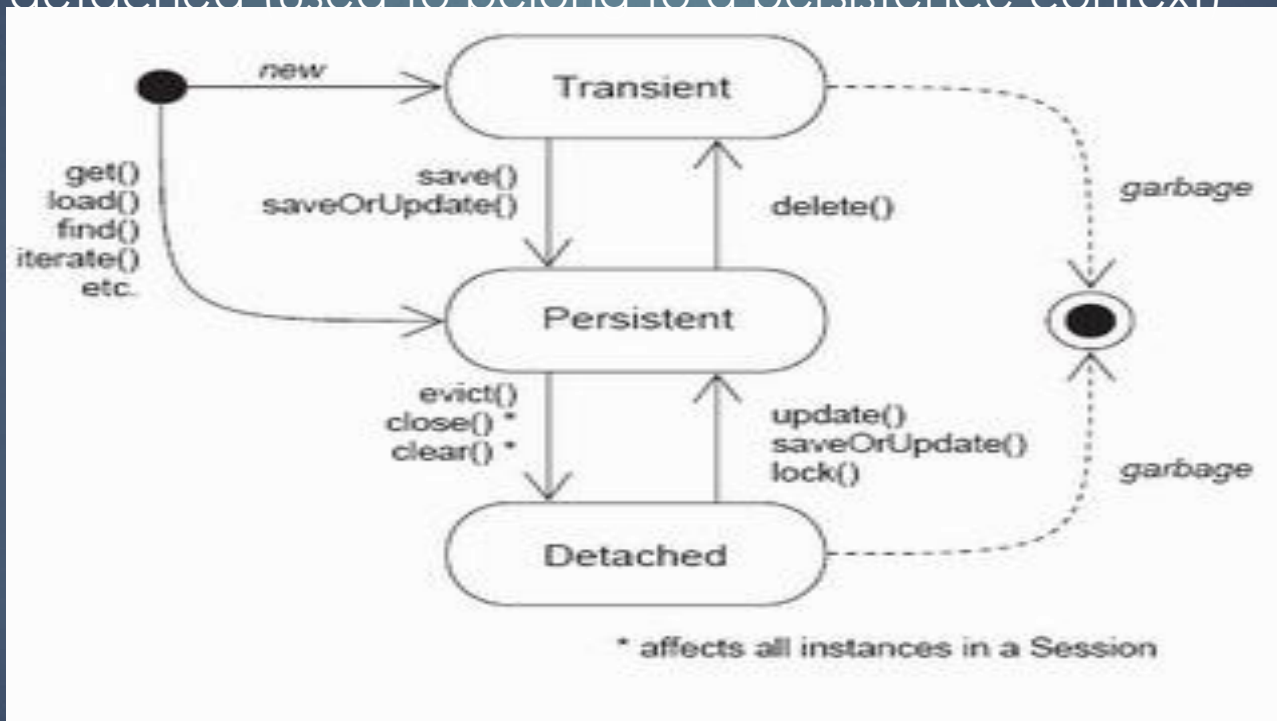
```
Session sess = sf.openSession();  
    Order o=new Order();    // search and return  
Query q = session.createQuery("from Order order where "+  
                                +"order.id=:id");  
  
q.setString("id",name);  
List result = q.list();  
if (list.size() == 0) {  
    System.out.println("No Order having id "  
                        + name);  
    System.exit(0);  
}  
o = (Order) list.get(0);  
sess.close();  
  
o.getOrderDate();
```

Persistent state

Detached State

Instance States

- ▶ An instance of a persistent classes may be in one of three different states, which are defined with respect to a persistence context
 - ▶ transient (does not belong to a persistence context)
 - ▶ persistent (belongs to a persistence context)
 - ▶ detached (used to belong to a persistence context)



Instance States

▶ “transient” state

- ▶ The instance is not, and has never been associated with any session (persistence context)
- ▶ It has no persistent identity (primary key value)
- ▶ It has no corresponding row in the database – ex) When POJO instance is created outside of a session

▶ “persistent” state

- ▶ The instance is currently associated with a session (persistence context).
- ▶ It has a persistent identity (primary key value) and likely to have a corresponding row in the database – ex) When an object is created within a session or a transient object gets persisted

Instance States

- ▶ **“detached”**
 - ▶ The instance was once associated with a persistence context, but that context was closed, or the instance was serialized to another process
 - ▶ It has a persistent identity and, perhaps, a corresponding row in the database
 - ▶ Used when POJO object instance needs to be sent over to another program for manipulation without having persistent context

State Transitions

- ▶ Transient instances may be made persistent by calling *save()*, *persist()* or *saveOrUpdate()*
- ▶ Persistent instances may be made transient by calling *delete()*
- ▶ Any instance returned by a *get()* or *load()* method is persistent
- ▶ Detached instances may be made persistent by calling *update()*, *saveOrUpdate()*, *lock()* or *replicate()*
- ▶ The state of a transient or detached instance may also be made persistent as a new persistent instance by calling *merge()*.

saveOrUpdate

- ▶ if the object is already persistent in this session, do nothing
- ▶ if another object associated with the session has the same identifier, throw an exception
- ▶ if the object has no identifier property, save() it
- ▶ if the object's identifier has the value assigned to a newly instantiated object, save() it

Persist

- ▶ Save and persist return type:
- ▶ Both INSERT records into database but **return type of persist is void** while return type of save is Serializable object.
- ▶ transient to instance persistent.
- ▶ persist() method doesn't guarantee that the identifier value will be assigned to the persistent instance immediately, the assignment might happen at flush time.
- ▶ behavior on outside of transaction boundaries.
- ▶ persist() guarantees that it will **not** execute an INSERT statement if it is called outside of transaction boundaries.
- ▶ save() method **does not** guarantee the same, if you are inside or outside of a transaction.

Clear()

- ▶ it explicitly is to remove all managed entities from L1 cache, so that it does not grow infinitely when processing large data sets in one transaction.
- ▶ It discards all the changes that were made to managed entites not explicitly persisted.
- ▶ This means that you can safely modify an entity, update it explicitly and clear the session.



Hibernate Mapping

Association -Relationship

- ▶ Association is a relationship of one class to another class
 - ▶ Its '**relationships**' in database terminology.
- ▶ Database design involves the master detail tables to represent the relationship of one entity to another.
- ▶ **A many-to-one Association**
 - ▶ Each book is related to one publisher
- ▶ **A one-to- Many Association**
 - ▶ One publisher may publish many books.
- ▶ **A One-To-One Association**
 - ▶ One Student has one Teacher
- ▶ **A Many-To-Many**
 - ▶ Many customers have many bank accounts
- ▶ **“unidirectional” association.**
 - ▶ If Association is navigable from book to publisher only

Simple Association (One-to-One)

- ▶ Expresses a relationship between two classes where each instance of the first class is related to a single instance of the second or vice versa
- ▶ Can be expressed in the database in two ways
 1. Giving each of the respective tables the same primary key values
 2. Using Foreign Key constraint from one table onto a unique identifier Column of the other.

One-to-Many-Relationship

- ▶ A one-to-many reference is basically a collection. Here class, OrderBook, holds a reference to a collection of another class, Order.
- ▶ In general <set> type collections are hibernate supports various types of collections
- ▶ We create an extra column in table Order which holds the FK to table OrderBook
- ▶ This allows OrderBook to be assigned a collection of Order based on the value of the cust_id column in ob_cust_id

1-N- “UD” Class and Table Design

OrderBook

```
private int cust_id;  
private String location;  
private Set orders;
```

Order

```
private int order_id ;  
private String customer_name;  
private double orderValue;  
private int ob_cust_id;;
```

OrderBook

cust_id, int(10) PRI
location, varchar(20)

Participant

order_id, int(10) PRI
orderValue, double
ob_cust_id, int(10) FK
customer_name, varchar(20)



Hibernate Query

Hibernate Query Language (HQL)

- Very similar to SQL but less verbose
- Understands OO – inheritance, polymorphism, association, aggregation, composition
- Selection: *from, as*
- Associations and joins: *inner join, outer join, right outer join, full join*
 - Projection: *select, elements*
 - Constraints: *where*
 - Other constructs: *aggregate functions, expressions, order by clauses, group by clauses, polymorphic selections, sub-queries*
- **Differences from SQL**
- HQL is fully object-oriented, understanding notions like inheritance, polymorphism and association
- Queries are case-insensitive, except for names of Java classes and properties

Hibernate Query Capabilities

- Criteria Queries
 - Extensible framework for expressing query criteria as objects
 - Includes “query by example”
- Native SQL Queries
- Enhanced support for queries expressed in the native SQL dialect of the database
 - “from” clause
 - Associations and join
 - “select” clause
 - Polymorphic query
 - “where” clause

Querying Objects

- The **list()** Method of the Query retrieves the result list containing objects
 - `Query query = session.createQuery("from Book");`
 - `List books = query.list();`
- To get One unique object returned as result,
 - **“?” to represent a query parameter and set it by index , which is zero-based not one-based as in JDBC.**

```
Query query = session.createQuery("from Book where isbn = ?");  
query.setString(0, "1932394419");  
Book book = (Book) query.uniqueResult();
```

```
Query query = session.createQuery("from Book where isbn =  
:isbn");  
query.setString("isbn", "1932394419");  
Book book = (Book) query.uniqueResult();
```

where clause

- The where clause allows you to narrow the list of instances returned.
- If no alias exists, you may refer to properties by name
- ```
Query qry = sess.createQuery("from Invoice where invNo=:ino");
 qry.setInteger("ino", 1001);
 Invoice inv = (Invoice)qry.uniqueResult();
```
- If there is an alias, use a qualified property name:

```
Query qry = sess.createQuery("from Book as bk where bk.id=?");

qry.setInteger(0, pubId);
```



# select clause

- Queries may return properties of any value type including properties of component type
  - `Query qry =mySess.createQuery("select invNo,customer from Invoice");`  
  
`List myList = qry.list();`  
  
`for(int i =0;i<myList.size();i++)`  
`{`  
`Object[] obj = (Object[]) myList.get(i);`  
  
`System.out.println(obj[0]);`  
`System.out.println(obj[1]);`  
`}`  
  
`Query qry =mySess.createQuery("select new Invoice(invNo,customer) from Invoice")`

## Aggregate functions

HQL queries may even return the results of aggregate functions on properties:

```
select avg(cat.weight), sum(cat.weight),
 max(cat.weight), count(cat)
from Cat cat
```

The supported aggregate functions are

- avg(...), sum(...), min(...), max(...)
- count(\*)
- count(...), count(distinct ...), count(all...)

You may use arithmetic operators, concatenation, and recognized SQL functions in the select clause:

```
select cat.weight + sum(kitten.weight)
from Cat cat
```

# Native SQL Query

- `String sql="select * from invoice_table where invNumber=101 and custName='Ramesh'";`

`Query qry= mySession.createQuery(sql).addEntity(Invoice.class);`

- To return a list of scalars or values as `Object[]` for each column.
- Hibernate uses `ResultSetMetadata` to infer column types can explicitly denote return types using `addScalar()`.

`ResultSetMetadata` to deduce the actual order and types

```
sess.createQuery("SELECT * FROM Customer").list();
```

```
sess.createQuery("SELECT ID, NAME, BIRTHDATE FROM
Customer").list();
```

# Native SQL Query

```
public List<Customer> scalarQuery() {
```

```
 Session sess=fact.openSession();
```

```
 Query qry =sess.createQuery
```

```
("Select customerNumber,CustomerName from Hiber_CUSTOMER ").
```

```
 addScalar("customerNumber",Hibernate.INTEGER).
```

```
 addScalar("customerName",Hibernate.STRING);
```

```
 List<Customer> custList =qry.list();
```

```
 sess.close();
```

```
 return custList;
```

```
}
```

```
for(int i =0;i<custList.size();i++) {
```

```
 Object[] val =(Object[]) custList.get(i);
```

```
 System.out.println(val[0]);
```

# Named Queries

- HQL statements in the mapping definitions are called “Named Queries”, The named queries can be put in any mapping definitions.
- For each named query, we need to assign a unique name to it. We should also put the query string in a `<![CDATA[...]]>` block to avoid conflicts with the special XML characters.

```
<hibernate-mapping package="com.training">
 <class name="Invoice" table="Invoice">
 - - - -
 </class>
 <query name="amountQry">
 <![CDATA[from Invoice where invAmount>?]]>
 </query>
</hibernate-mapping>
```

```
Session sess = HibernateUtil.getSession();
```

```
Query qry = sess.getNamedQuery("amountQry");
qry.setDouble(0, 1000);
```

```
List<Invoice> invList = qry.list();
```

# Using Stored Procedure

```
CREATE PROCEDURE `GetStocks`(int_stockcode VARCHAR(20))
BEGIN
 SELECT * FROM stock WHERE stock_code = int_stockcode;
 END $$

DELIMITER ;
```

```
Query query = session.createSQLQuery("CALL GetStocks(:stockCode)")
 .addEntity(Stock.class) .setParameter("stockCode", "7277");

List result = query.list();
for(int i=0; i<result.size(); i++) {
 Stock stock = (Stock)result.get(i);
 System.out.println(stock.getStockCode());
}
```



# Hibernate Annotations

- ▶ **@Entity** - class level
  - ▶ POJO class can be declared as an entity
- ▶ **@Id** - Property Level or at all the getXXX Level
  - ▶ declares the identifier property of the entity.
- ▶ **@Table** - class level
  - ▶ to define the table, catalog, and schema names for entity mapping.
- ▶ **@Table**(name="tbl\_name",  
uniqueConstraints =

# Hibernate Annotations

```
import javax.persistence.*;
```

```
@Entity
```

```
@Table(name = "employee")
```

```
public class Employee {
```

```
 @Id
```

```
 @GeneratedValue(strategy=GenerationType.AUTO)
```

```
 @Column(name = "id")
```

```
 Integer id;
```

# Using Annotations

- ▶ In the Hibernate configuration File
- ▶ **<mapping class="com.training.CreditCard"/>**
- ▶ In the Application Class
- ▶ **SessionFactory session = new  
AnnotationConfiguration().configure().buildSessionFactory();**